
Cardiovascular Disease Classification

Chance, Dan, David and Shane

May 23, 2020

TABLE OF CONTENTS:

1	Business Understanding	1
1.1	Library Imports	1
1.2	Load the CSV Data	1
2	Data Meaning Type	3
3	Data Quality	5
3.1	Age	7
3.2	Height	7
3.3	Weight	8
3.4	Systolic blood pressure (ap_hi)	9
3.5	Diastolic blood pressure (ap_lo)	10
3.5.1	Categorical Variables	11
3.6	Age	14
3.7	Height	15
3.8	Weight	15
3.9	Systolic blood pressure (ap_hi)	16
3.10	Diastolic blood pressure (ap_lo)	17
4	Simple Statistics	19
5	Visualize Attributes	23
6	Explore Joint Attributes	27
7	Explore Attributes and Class	31
8	New Features	39
9	Exceptional Work	47
9.1	Store feature matrix as an ndarray	47
9.2	Store response vector	47

BUSINESS UNDERSTANDING

The provided cardiovascular dataset is available for academic purposes, aimed at classifying whether a patient may have cardiovascular disease. Stated differently, this data was collected and created to build and train models using a variety of predictors to classify a patient as having cardiovascular disease or not having cardiovascular disease. The outcomes of the cardiovascular dataset can be defined as a classification problem of a categorical response variable. For example, to determine whether smoking is a predictor of cardiovascular disease and using accuracy statistics to measure the effectiveness of smoking as a predictor. If the model correctly predicts the presence of cardiovascular disease in a given a patient 70% of the time actions such as early, noninvasive, and safe intervention therapies (such as with diet and exercise), may replace necessary (surgical, and dangerous interventions) as an undiagnosed problem worsens. A low false-negative rate, that is incorrectly diagnosing a patient as not having cardiovascular disease may be a beneficial measure of having mined useful knowledge. Given the categorical response variable and classification of cardiovascular disease, the effectiveness of the prediction algorithm may be measured through accuracy or precision statistics, ROC curves or cross-validation (“CV”), for example a 10-fold CV.

1.1 Library Imports

```
[1]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
from matplotlib.ticker import MaxNLocator
import seaborn as sns
import plotly.express as px
%matplotlib inline
```

1.2 Load the CSV Data

```
[2]: df = pd.read_csv('../data/cardio_train.csv', delimiter=';')
```


DATA MEANING TYPE

The cardiovascular dataset consists of 11 features and 1 target variable. Attribute descriptions can be found on [Kaggle](#). They are provided below.

There are 3 types of input features in the data set. - **Objective**: factual information - **Examination**: results from a medical examination - **Subjective**: response provided by patient

Table 1: Cardiovascular Dataset - Attribute Descriptions

Column Description	Feature Type	Column Name	Data Type
Age	Objective	age	int (days)
Height	Objective	height	int (cm)
Weight	Objective	weight	float (kg)
Gender	Objective	gender	1: female, 2: male
Systolic blood pressure	Examination	ap_hi	int
Diastolic blood pressure	Examination	ap_lo	int
Cholesterol	Examination	cholesterol	1: normal, 2: above normal, 3: well above normal
Glucose	Examination	gluc	1: normal, 2: above normal, 3: well above normal
Smoking	Subjective	smoke	binary
Alcohol intake	Subjective	alco	binary
Physical activity	Subjective	active	binary
Has CVD?	Target	cardio	binary

```
[3]: df.head()
```

```
[3]:
```

	id	age	gender	height	weight	ap_hi	ap_lo	cholesterol	gluc	smoke	\
0	0	18393	2	168	62.0	110	80	1	1	0	
1	1	20228	1	156	85.0	140	90	3	1	0	
2	2	18857	1	165	64.0	130	70	3	1	0	
3	3	17623	2	169	82.0	150	100	1	1	0	
4	4	17474	1	156	56.0	100	60	1	1	0	

	alco	active	cardio
0	0	1	0
1	0	1	1
2	0	0	1
3	0	1	1
4	0	0	0

We decided to use the `id` column as the index as it appears to be a unique identifier for the subject.

```
[4]: df.set_index("id", inplace=True)
```

```
[5]: df.index.is_unique
```

```
[5]: True
```


DATA QUALITY

```
[6]: df.shape
[6]: (70000, 12)
```

Our data set has 70 thousand rows and 12 columns, which should work nicely for many of the machine learning classifiers we may attempt to utilize.

```
[7]: cols = df.columns
```

Store columns for later use.

```
[8]: cols
[8]: Index(['age', 'gender', 'height', 'weight', 'ap_hi', 'ap_lo', 'cholesterol',
          'gluc', 'smoke', 'alco', 'active', 'cardio'],
          dtype='object')
```

```
[9]: df.isna().any()
[9]: age           False
     gender        False
     height        False
     weight        False
     ap_hi         False
     ap_lo         False
     cholesterol   False
     gluc          False
     smoke         False
     alco          False
     active        False
     cardio        False
     dtype: bool
```

There are no missing values in any of the columns of the default data. But upon further inspection we found that zeros or other values may have been used in place of missing entries which should also be addressed.

```
[10]: df.duplicated().any()
[10]: True
```

There do appear to be missing values.

```
[11]: df.duplicated().sum()
```

```
[11]: 24
```

There were a total of 24 sets where all columns of an observation equal at least one other record in the set. It was important to remove `id` as part of the data frame otherwise these duplicated entries would have been more difficult to detect.

```
[12]: duplicated = df[df.duplicated(keep=False)].sort_values(by=['age', 'gender', 'height',
↳ 'weight', 'ap_hi', 'ap_lo', 'cholesterol',
    'gluc', 'smoke', 'alco', 'active', 'cardio'])
```

```
[13]: duplicated.head(10)
```

```
[13]:
```

	age	gender	height	weight	ap_hi	ap_lo	cholesterol	gluc	smoke	\
id										
9004	14552	1	158	64.0	120	80	1	1	0	
57690	14552	1	158	64.0	120	80	1	1	0	
24435	16160	1	168	65.0	120	80	1	1	0	
91592	16160	1	168	65.0	120	80	1	1	0	
1685	16793	1	165	68.0	120	80	1	1	0	
31110	16793	1	165	68.0	120	80	1	1	0	
40450	16805	1	157	67.0	120	80	1	1	0	
86345	16805	1	157	67.0	120	80	1	1	0	
14974	16937	2	170	70.0	120	80	1	1	0	
63776	16937	2	170	70.0	120	80	1	1	0	

	alco	active	cardio
id			
9004	0	1	0
57690	0	1	0
24435	0	1	1
91592	0	1	1
1685	0	1	0
31110	0	1	0
40450	0	1	0
86345	0	1	0
14974	0	0	0
63776	0	0	0

```
[14]: df_clean = df.copy(deep=True)
```

```
[15]: df_clean.drop_duplicates(inplace=True)
```

We'll remove the duplicates entirely, as this should not affect the ability of our models to make predictions with the amount of observations at our disposal.

```
[16]: # %%time
# df_clean['age'] = df_clean['age'].apply(lambda x: round(x / 365))
df_clean['age'] = (df_clean['age'] / 365).round().astype('int')
```

Age was provided in days, and for the sake of interpretability we'll be converting this to years for all observations.

```
[17]: plt.style.use('ggplot')
```

3.1 Age


```
[18]: fig_1 = plt.figure(1, figsize=(20, 5))

chart_1 = fig_1.add_subplot(121)
chart_2 = fig_1.add_subplot(122)

chart_1.hist(df_clean["age"])
chart_1.xaxis.set_major_locator(MaxNLocator(integer=True))
chart_1.set_title('Histogram of Age')
chart_1.set_xlabel('Age (in years)')
chart_1.set_ylabel('Frequency')

sns.boxplot(x="age", data=df_clean, ax=chart_2)
chart_2.xaxis.set_major_locator(MaxNLocator(integer=True))
chart_2.set_title('Boxplot of Age')
chart_2.set_xlabel('Age (in years)')

plt.show()
```



Age has relatively few outliers but is slightly right-skewed.

3.2 Height

```
[19]: fig_2 = plt.figure(1, figsize=(20, 5))

chart_1 = fig_2.add_subplot(121)
chart_2 = fig_2.add_subplot(122)


chart_1.hist(df_clean["height"])
chart_1.xaxis.set_major_locator(MaxNLocator(integer=True))
chart_1.set_title('Histogram of Height')
chart_1.set_xlabel('Height (in cm)')
chart_1.set_ylabel('Frequency')

sns.boxplot(x="height", data=df_clean, ax=chart_2)
chart_2.xaxis.set_major_locator(MaxNLocator(integer=True))
chart_2.set_title('Boxplot of Height')
chart_2.set_xlabel('Height (in cm)')
```

(continues on next page)

(continued from previous page)

```
plt.show()
```



There are quite a few outliers in the height column that should be addressed. The largest of 250cm is over 8 feet tall and appears to be an error.

3.3 Weight

```
[20]:
```


```
fig_3 = plt.figure(1, figsize=(20, 5))

chart_1 = fig_3.add_subplot(121)
chart_2 = fig_3.add_subplot(122)

chart_1.hist(df_clean["weight"])
chart_1.xaxis.set_major_locator(MaxNLocator(integer=True))
chart_1.set_title('Histogram of Weight')
chart_1.set_xlabel('Weight (in kg)')
chart_1.set_ylabel('Frequency')

sns.boxplot(x="weight", data=df_clean, ax=chart_2)
chart_2.xaxis.set_major_locator(MaxNLocator(integer=True))
chart_2.set_title('Boxplot of Weight')
chart_2.set_xlabel('Weight (in kg)')

plt.show()
```



Similarly, there are a lot of outliers in the weight column as well.

3.4 Systolic blood pressure (ap_hi)


```
[21]: fig_4 = plt.figure(1, figsize=(20, 5))

chart_1 = fig_4.add_subplot(121)
chart_2 = fig_4.add_subplot(122)

chart_1.hist(df_clean["ap_hi"])
chart_1.xaxis.set_major_locator(MaxNLocator(integer=True))
chart_1.set_title('Histogram of Systolic blood pressure')
chart_1.set_xlabel('ap_hi')
chart_1.set_ylabel('Frequency')

sns.boxplot(x="ap_hi", data=df_clean, ax=chart_2)
chart_2.xaxis.set_major_locator(MaxNLocator(integer=True))
chart_2.set_title('Boxplot of Systolic blood pressure')
chart_2.set_xlabel('ap_hi')

plt.show()
```



The distribution of the Systolic blood pressure was quite unusual with several readings that were likely erroneous.

```
[22]: df_clean["ap_hi"].sample(10)
```

```
[22]: id
13160    130
29293    115
61102    150
90710    120
72912    120
90865    120
63998    140
78175    120
90864    120
96667    110
Name: ap_hi, dtype: int64
```

A random sample show values within the expected range.

```
[23]: df_clean["ap_hi"].sort_values()

[23]: id
50055      -150
34295      -140
66571      -120
36025      -120
22881      -115
...
36339     14020
66998     14020
67502     14020
36414     14020
58374     16020
Name: ap_hi, Length: 69976, dtype: int64
```

But there were negative values and extremely high ones that should be reviewed. We'll address these outliers later within the imputation section.

```
[24]: # df_clean = df_clean[~(df_clean['ap_hi'] < 40) & (df_clean['ap_hi'] < 300)]
# df_clean.shape[0]
```

3.5 Diastolic blood pressure (ap_lo)

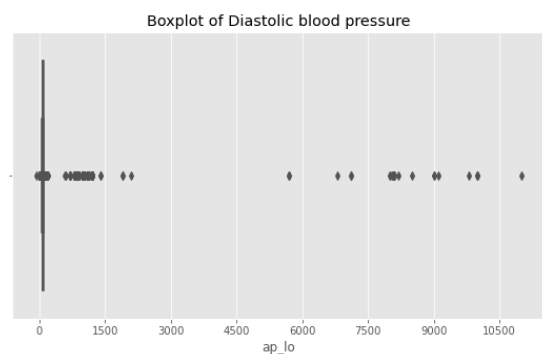

```
[25]: fig_5 = plt.figure(1, figsize=(20, 5))

chart_1 = fig_5.add_subplot(121)
chart_2 = fig_5.add_subplot(122)

chart_1.hist(df_clean["ap_lo"])
chart_1.xaxis.set_major_locator(MaxNLocator(integer=True))
chart_1.set_title('Histogram of Diastolic blood pressure')
chart_1.set_xlabel('ap_lo')
chart_1.set_ylabel('Frequency')

sns.boxplot(x="ap_lo", data=df_clean, ax=chart_2)
chart_2.xaxis.set_major_locator(MaxNLocator(integer=True))
chart_2.set_title('Boxplot of Diastolic blood pressure')
chart_2.set_xlabel('ap_lo')

plt.show()
```



```
[26]: df_clean["ap_lo"].sample(10)
```

```
[26]: id
68416      80
73487      80
72381     100
93163      70
37294      80
89255      70
43796      70
61619      70
90459      80
47189      90
Name: ap_lo, dtype: int64
```

```
[27]: df_clean["ap_lo"].sort_values()
```

```
[27]: id
85816      -70
98095       0
45400       0
75007       0
81298       0
...
62058     9800
34098    10000
3352     10000
97907    10000
61901    11000
Name: ap_lo, Length: 69976, dtype: int64
```

The same technique should be applied to the `ap_lo` feature.

3.5.1 Categorical Variables


```
[28]: fig_0 = plt.figure(1, figsize=(20, 5))

chart_1 = fig_0.add_subplot(121)
chart_2 = fig_0.add_subplot(122)

sns.countplot(x="cardio", hue="cardio", data=df_clean, ax=chart_1)
# chart_1.legend(bbox_to_anchor=(1,1), title='CVD')
chart_1.set_title('Cardiovascular Disease')
chart_1.set_xlabel('CVD')
chart_1.set_ylabel('Count')

pd.crosstab(df_clean["cholesterol"], df_clean["cardio"]).plot(kind="bar", ax=chart_2)
chart_2.set_title('CVD by Cholesterol Level')
chart_2.set_xlabel('Cholesterol Level')
chart_2.set_ylabel('Count')

plt.show()
```



The Cardiovascular Disease (CVD) response variable is equally distributed. The presence or absence of CVD does seem to change with the cholesterol levels.

```
[29]: fig_6 = plt.figure(1, figsize=(20, 5))

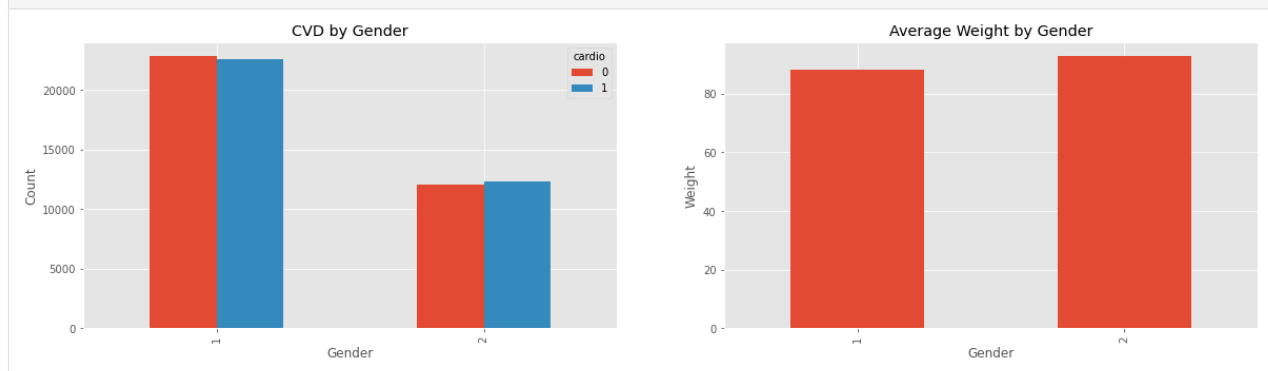
chart_1 = fig_6.add_subplot(121)
chart_2 = fig_6.add_subplot(122)

pd.crosstab(df_clean["gender"], df_clean["cardio"]).plot(kind="bar", ax=chart_1)
chart_1.set_title('CVD by Gender')
chart_1.set_xlabel('Gender')
chart_1.set_ylabel('Count')

df_clean.groupby(['gender', 'height']).sum().reset_index()

df_clean.groupby(['gender', 'weight']).mean().reset_index().groupby('gender')['weight
→'].mean().plot(kind='bar', ax=chart_2)
chart_2.set_title('Average Weight by Gender')
chart_2.set_xlabel('Gender')
chart_2.set_ylabel('Weight')

plt.show()
```



There are more subjects with the label 1 in the study than those with label 2.

We're going to assume that label 2 is for male as the mean weight is slightly heigher for that category.

```
[30]: fig_7 = plt.figure(1, figsize=(20, 5))

chart_1 = fig_7.add_subplot(121)
```

(continues on next page)

(continued from previous page)

```


chart_2 = fig_7.add_subplot(122)

pd.crosstab(df_clean["gluc"], df_clean["cardio"]).plot(kind="bar", ax=chart_1)
chart_1.set_title('CVD by Glucose Level')
chart_1.set_xlabel('Glucose Level')
chart_1.set_ylabel('Count')

pd.crosstab(df_clean["smoke"], df_clean["cardio"]).plot(kind="bar", ax=chart_2)
chart_2.set_title('CVD by Smoking Level')
chart_2.set_xlabel('Smoking Level')
chart_2.set_ylabel('Count')

plt.show()

```



The presence or absence of CVD also changes with the glucose levels but suprsingly not with smoking.

```
[31]: df_clean.groupby(['gluc', 'cardio']).size().unstack(fill_value=0)
```

```

[31]: cardio      0      1
      gluc
      1      30877  28578
      2       2112   3078
      3       2015   3316

```

There are only a few thousand entries within levels 2 and 3 of the Choleserol column.

```

[32]: fig_8 = plt.figure(1, figsize=(20, 5))

      chart_1 = fig_8.add_subplot(121)
      chart_2 = fig_8.add_subplot(122)

      pd.crosstab(df_clean["alco"], df_clean["cardio"]).plot(kind="bar", ax=chart_1)
      chart_1.set_title('CVD by Alchohol Level')
      chart_1.set_xlabel('Alcohol Level')
      chart_1.set_ylabel('Count')


      pd.crosstab(df_clean["active"], df_clean["cardio"]).plot(kind="bar", ax=chart_2)
      chart_2.set_title('CVD by Activity Level')
      chart_2.set_xlabel('Activity Level')
      chart_2.set_ylabel('Count')

```

(continues on next page)

(continued from previous page)

plt.show()



Also surprising is that the Alcohol level didn't seem to have an impact on the response variable. The Activity Level did show good separation.

Imputation

3.6 Age

```
[33]: median_age = df_clean['age'].median()
```

```
[34]: age_outlier_ids = df_clean.index[(np.abs(df_clean['age'] - df_clean['age'].mean()) >
↳ (3 * df_clean['age'].std()))]
```

```
[35]: df_clean.loc[df_clean.index.isin(age_outlier_ids), "age"] = median_age
```

```
[36]: df_clean.loc[df_clean.index.isin(age_outlier_ids)].head()
```

```
[36]:
```

	age	gender	height	weight	ap_hi	ap_lo	cholesterol	gluc	smoke	\
id										
8850	54.0	1	175	59.0	120	80	1	1	0	
31922	54.0	2	175	92.0	100	60	1	1	0	
43842	54.0	1	159	59.0	120	80	1	1	0	
79749	54.0	1	160	59.0	110	70	1	1	0	

	alco	active	cardio
id			
8850	0	1	0
31922	0	1	0
43842	0	1	0
79749	0	1	0

We've imputed observations with an age of more than three standard deviations from the mean with the median value. (~4 observations in total) These may well have been valid observations but we wanted our model to extend well to other new and unseen data sets.

3.7 Height

```
[37]: median_height = df_clean['height'].median()

[38]: height_outlier_ids = df_clean.index[(np.abs(df_clean['height'] - df_clean['height'].
↳mean()) > (3 * df_clean['height'].std()))]

[39]: df_clean.loc[df_clean.index.isin(height_outlier_ids), "height"] = median_height

[40]: df_clean.loc[df_clean.index.isin(height_outlier_ids)].head()

[40]:
```

	age	gender	height	weight	ap_hi	ap_lo	cholesterol	gluc	smoke	\
id										
309	60.0	2	165.0	55.0	120	80	1	1	0	
1432	42.0	2	165.0	111.0	148	86	1	1	0	
1554	61.0	2	165.0	68.0	120	80	1	1	0	
3049	44.0	2	165.0	74.0	140	90	1	1	1	
3399	62.0	2	165.0	52.0	140	100	1	1	0	

	alco	active	cardio
id			
309	0	1	0
1432	0	1	1
1554	0	1	1
3049	1	1	1
3399	0	1	1

We've imputed observations with a height more than three standard deviations from the mean with the median value. (~287 observations in total)

3.8 Weight

```
[41]: median_weight = df_clean['weight'].median()

[42]: weight_outlier_ids = df_clean.index[(np.abs(df_clean['weight'] - df_clean['weight'].
↳mean()) > (3 * df_clean['weight'].std()))]

[43]: df_clean.loc[df_clean.index.isin(weight_outlier_ids), "weight"] = median_weight

[44]: df_clean.loc[df_clean.index.isin(weight_outlier_ids)].head()

[44]:
```

	age	gender	height	weight	ap_hi	ap_lo	cholesterol	gluc	smoke	\
id										
474	58.0	1	157.0	72.0	120	80	1	1	0	
552	46.0	2	165.0	72.0	120	80	1	1	0	
618	46.0	1	186.0	72.0	130	70	1	1	0	
634	58.0	2	178.0	72.0	160	90	1	3	0	
846	40.0	1	164.0	72.0	130	90	1	1	0	

	alco	active	cardio
id			
474	0	1	0
552	0	1	0
618	0	1	0
634	1	1	0
846	0	1	0

(continues on next page)

(continued from previous page)

id			
474	0	1	1
552	0	1	0
618	0	0	0
634	0	1	1
846	0	0	0

The same method was applied to the 702 weight outliers.

3.9 Systolic blood pressure (ap_hi)

```
[45]: median_ap_hi = df_clean['ap_hi'].median()
```

```
[46]: ap_hi_outlier_ids = df_clean.index[(np.abs(df_clean['ap_hi'] - df_clean['ap_hi'].
↳mean()) > (3 * df_clean['ap_hi'].std()))]
```

```
[47]: df_clean.loc[df_clean.index.isin(ap_hi_outlier_ids), "ap_hi"] = median_ap_hi
```

```
[48]: df_clean.loc[df_clean.index.isin(ap_hi_outlier_ids)].head()
```

```
[48]:
```

	age	gender	height	weight	ap_hi	ap_lo	cholesterol	gluc	smoke	\
id										
2654	41.0	1	160.0	60.0	120.0	60	1	1	0	
2845	62.0	2	167.0	59.0	120.0	0	1	1	0	
6822	40.0	1	168.0	63.0	120.0	60	2	1	0	
11089	58.0	1	175.0	80.0	120.0	90	1	1	0	
12710	52.0	1	164.0	75.0	120.0	80	2	1	0	

	alco	active	cardio
id			
2654	0	1	0
2845	0	1	0
6822	0	1	0
11089	0	1	1
12710	0	1	1

```
[49]: ap_hi_outlier_ids2 = df_clean.index[(df_clean['ap_hi'] < 40) | (df_clean['ap_hi'] >_
↳300)]
```

```
[50]: df_clean.loc[df_clean.index.isin(ap_hi_outlier_ids2), "ap_hi"] = median_ap_hi
```

Even after handling the ap_hi outliers through the standard deviation method, there were still some unusual entries that were manually addressed. (readings less than 40 or greater than 300)

3.10 Diastolic blood pressure (ap_lo)

```
[51]: median_ap_lo = df_clean['ap_lo'].median()
```

```
[52]: ap_lo_outlier_ids = df_clean.index[(np.abs(df_clean['ap_lo'] - df_clean['ap_lo'].
↪mean()) > (3 * df_clean['ap_lo'].std()))]
```

```
[53]: df_clean.loc[df_clean.index.isin(ap_lo_outlier_ids), "ap_lo"] = median_ap_lo
```

```
[54]: df_clean.loc[df_clean.index.isin(ap_lo_outlier_ids)].head()
```

```
[54]:
```

	age	gender	height	weight	ap_hi	ap_lo	cholesterol	gluc	smoke	\
id										
314	48.0	2	183.0	98.0	160.0	80.0	1	2	1	
334	60.0	2	157.0	60.0	160.0	80.0	2	1	0	
357	50.0	1	150.0	83.0	140.0	80.0	1	1	0	
458	64.0	1	176.0	63.0	160.0	80.0	2	2	0	
482	51.0	1	154.0	81.0	140.0	80.0	2	1	0	

	alco	active	cardio
id			
314	0	1	1
334	0	0	1
357	0	1	1
458	0	0	1
482	0	1	1

```
[55]: ap_lo_outlier_ids2 = df_clean.index[(df_clean['ap_lo'] < 40) | (df_clean['ap_lo'] >_
↪300)]
```

```
[56]: df_clean.loc[df_clean.index.isin(ap_lo_outlier_ids2), "ap_lo"] = median_ap_lo
```

The ap_lo feature needed similar processing.

SIMPLE STATISTICS

```
[57]: df_clean.describe()
```

```
[57]:
```

	age	gender	height	weight	ap_hi \
count	69976.000000	69976.000000	69976.000000	69976.000000	69976.000000
mean	53.340317	1.349648	164.411884	73.643909	126.981551
std	6.763333	0.476862	7.708061	13.193980	17.075701
min	39.000000	1.000000	140.000000	32.000000	60.000000
25%	48.000000	1.000000	159.000000	65.000000	120.000000
50%	54.000000	1.000000	165.000000	72.000000	120.000000
75%	58.000000	2.000000	170.000000	81.000000	140.000000
max	65.000000	2.000000	188.000000	117.000000	240.000000

	ap_lo	cholesterol	gluc	smoke	alco \
count	69976.000000	69976.000000	69976.000000	69976.000000	69976.000000
mean	81.365125	1.366997	1.226535	0.088159	0.053790
std	9.624691	0.680333	0.572353	0.283528	0.225604
min	40.000000	1.000000	1.000000	0.000000	0.000000
25%	80.000000	1.000000	1.000000	0.000000	0.000000
50%	80.000000	1.000000	1.000000	0.000000	0.000000
75%	90.000000	2.000000	1.000000	0.000000	0.000000
max	190.000000	3.000000	3.000000	1.000000	1.000000

	active	cardio
count	69976.000000	69976.000000
mean	0.803718	0.499771
std	0.397187	0.500004
min	0.000000	0.000000
25%	1.000000	0.000000
50%	1.000000	0.000000
75%	1.000000	1.000000
max	1.000000	1.000000

In the table above we have an overview of all of our attributes included in the dataset. This gives an idea of our counts after the cleaning of our data and provides us with data we can hope to draw inferences from. Below we will look at the simplest level of some of our attributes to review if some general assumptions are true or false within our given data.

```
[58]: df_clean.agg({'age': ['min', 'max', 'median', 'skew', 'std'], 'gender': ['min', 'max', 'median',
→ 'skew', 'std'], 'height': ['min', 'max', 'median', 'skew', 'std'], 'weight': ['min', 'max',
→ 'median', 'skew', 'std'], 'ap_hi': ['min', 'max', 'median', 'skew', 'std'], 'ap_lo': ['min',
→ 'max', 'median', 'skew', 'std']})
```

```
[58]:
```

	age	gender	height	weight	ap_hi	ap_lo
min	39.000000	1.000000	140.000000	32.000000	60.000000	40.000000
max	65.000000	2.000000	188.000000	117.000000	240.000000	190.000000
median	54.000000	1.000000	165.000000	72.000000	120.000000	80.000000
skew	-0.305042	0.630605	0.088790	0.550598	0.918609	0.680160
std	6.763333	0.476862	7.708061	13.193980	17.075701	9.624691

This table narrows down our quantitative variables, giving us a quick look at the balance of the dataset. The variable “skew” helps us see in which way our observations lean, none of which were too far out of preferred parameters. We noticed that age of our population is slightly younger (leaning towards our min of 39), gender favors men slightly (as it is a higher value and 2 represents male), height is almost indiscernible, weight favors our max slightly, and blood pressure is slightly higher in our dataset. As stated previously, with the clean data we feel comfortable making assumptions from this data and have decided to move forward.

```
[59]: df_clean.groupby("cholesterol").mean()
```

```
[59]:
```

	age	gender	height	weight	ap_hi	ap_lo
cholesterol						
1	52.831879	1.359275	164.653272	72.626377	125.074903	
2	53.712326	1.329459	163.912452	75.766443	131.291444	
3	56.200471	1.311059	163.436152	77.736511	134.256385	

	ap_lo	gluc	smoke	alco	active	cardio
cholesterol						
1	80.524875	1.099139	0.085235	0.048089	0.802353	0.440175
2	83.085140	1.335218	0.105666	0.076134	0.799037	0.602157
3	84.783412	1.924870	0.086412	0.064344	0.818125	0.765435

Cholesterol has been a topic of debate for years given advancements in technology. We felt that this would possibly give us some insight into the weight that it would have in regards to classifying cardiovascular disease. This attribute classifies 1 as normal levels of cholesterol, 2 as being above normal, and 3 being well above normal

```
[60]: df_clean.groupby("smoke").mean()
```

```
[60]:
```

	age	gender	height	weight	ap_hi	ap_lo
smoke						
0	53.441221	1.299513	163.941166	73.374581	126.832824	81.293510
1	52.296645	1.868212	169.280597	76.429616	128.519857	82.105852

	cholesterol	gluc	alco	active	cardio
smoke					
0	1.364819	1.227389	0.029934	0.800523	0.502186
1	1.389528	1.217701	0.300535	0.836764	0.474793

As made clear from numerous medical studies, smoking can greatly affect health. In this attribute we hold that 0 is a non-smoker and 1 is a smoker. We observed here that height was higher in the smoking population (which could be attributed to more male smokers than females). Weight was slightly higher which could also correspond with the skewed height, but it is also known that smokers do generally weigh more than non-smokers from previous medical studies. What was most expected is that the overall blood pressure (both diastolic and systolic) were both higher than non-smokers.

```
[61]: df_clean.groupby("active").mean()
```

```
[61]:
```

	age	gender	height	weight	ap_hi	ap_lo
active						
0	53.473244	1.343939	164.542774	74.082461	126.954277	81.372552
1	53.307854	1.351043	164.379919	73.536808	126.988211	81.363311

(continues on next page)

(continued from previous page)

	cholesterol	gluc	smoke	alco	cardio
active					
0	1.353331	1.234365	0.073316	0.042155	0.535857
1	1.370335	1.224623	0.091784	0.056631	0.490959

Lastly, we wanted to evaluate the physical activity of our population. This attribute's levels are 0 for non-active and 1 for active. What was curious about this attribute is that the age is not skewed towards either older nor younger. Only ever so slightly is the age skewed towards younger persons, but generally physical activity is not determined by age through this dataset, where normally we'd expect as a population ages, they become less active. Weight was another variable we felt we'd expect to see a difference between active and non-active. Here we did find this to be true as those who are active were approximately a half kilogram lighter than non-active. This wasn't as great as we would have expected, but it does prove our general assumption. Additionally, we'd really expect the systolic and diastolic blood pressures to be affected by this variable, however there was notable nearly no change here as well. From this table, we'd draw a conclusion that physical activity does not necessarily help prevent or predict cardiovascular disease, but much more investigation is required to prove or disprove this entirely as well.

VISUALIZE ATTRIBUTES

```
[62]: fig = px.scatter_3d(df_clean.sample(200), x='ap_hi', y='ap_lo', z='age',
                        color='cardio')

fig.update_layout(margin=dict(l=0, r=0, b=0, t=0))

fig.show()
```

Data type cannot be displayed: application/vnd.plotly.v1+json


```
[63]: import base64, io, IPython
      from PIL import Image as PILImage

image = PILImage.open("../img/3d_scatterplot.png")

output = io.BytesIO()
image.save(output, format='PNG')
encoded_string = base64.b64encode(output.getvalue()).decode()

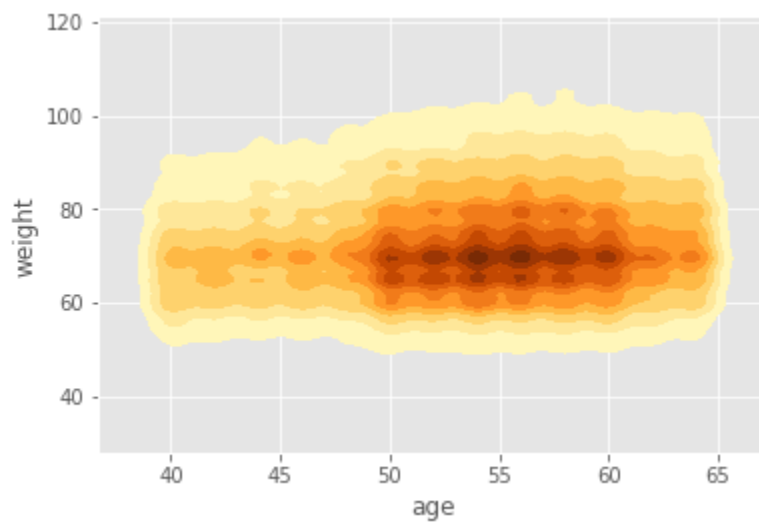
html = '<img src="data:image/png;base64,{}/>'.format(encoded_string)
IPython.display.HTML(html)
```

```
[63]: <IPython.core.display.HTML object>
```



The 3D scatterplot shows higher CVD response rates from ap_hi and ap_lo but not for age.

```
[65]: df_sub3 = df_clean[['age', 'weight', 'cholesterol', 'ap_hi', 'cardio', 'ap_lo', 'smoke', 'active']]
df_sub3.cardio = df_sub3.cardio=='1'
#Adapted from https://towardsdatascience.com/the-art-of-effective-visualization-of-multi-dimensional-data-6c7202990c57
ax = sns.kdeplot(df_sub3['age'], df_sub3['weight'],
                  cmap="YlOrBr", shade=True, shade_lowest=False)
```




The above Kernel Density plot depicts the density of those with cardiovascular disease by age and weight. As can

be seen, the most dense areas are shown with an age value of 54-56 and a weight value of about 70. This may be indicative of ages and weights that may be predictive of cardiovascular disease.

EXPLORE JOINT ATTRIBUTES

We can create a feature correlation heatmap using Seaborn.

```
[66]: features = ['age', 'gender', 'height', 'weight', 'ap_hi', 'ap_lo', 'cholesterol',  
↪ 'gluc', 'smoke', 'alco', 'active', 'cardio']  
  
plt.figure(figsize=(15,15))  
  
# Use an easier to see colormap  
cmap = sns.diverging_palette(220, 10, as_cmap=True)  
  
# Mask  
correlation = df_clean[features].corr()  
correlation[np.abs(correlation)<.2] = 0  
  
sns.heatmap(correlation, annot = True, cmap=cmap).set(title = 'Correlation Heatmap')  
plt.show()
```



The above heatmap shows some interesting correlations, namely between

- Height and Gender
- Gender and Smoking
- Alcohol and Smoking
- Glucose and Cholesterol

Additionally, we see the following are correlated with our target variable, 'cardio'.

- Age

- Blood Pressure (High/Low)
- Cholesterol

Let's cross tabulate a few of these and inspect further.

```
[67]: fig = plt.figure(1, figsize=(20, 5))


chart_1 = fig.add_subplot(131)
chart_2 = fig.add_subplot(132)
chart_3 = fig.add_subplot(133)

pd.crosstab(df_clean["gluc"], df_clean["cholesterol"]).apply(lambda r: r/r.sum(),
    ↪axis=1).plot(kind="bar", ax=chart_1)
chart_1.set_title('Glucose vs Cholesterol')
chart_1.set_xlabel('Glucose')
chart_1.set_ylabel('Percentage')

pd.crosstab(df_clean["alco"], df_clean["smoke"]).apply(lambda r: r/r.sum(), axis=1).
    ↪plot(kind="bar", ax=chart_2)
chart_2.set_title('Alcohol vs Smoking')
chart_2.set_xlabel('Alcohol')
chart_2.set_ylabel('Percentage')

pd.crosstab(df_clean["gender"], df_clean["smoke"]).apply(lambda r: r/r.sum(), axis=1).
    ↪plot(kind="bar", ax=chart_3)
chart_3.set_title('Gender vs Smoking')
chart_3.set_xlabel('Gender')
chart_3.set_ylabel('Percentage')

plt.show()
```



Some observations based on the above charts:

- We see above that glucose and cholesterol trend nicely. That is, if you are a x in one group, you are likely to be the same value in the other group.
- If you are a drinker, you are more than likely a smoker too.
- If you are a male, you are more likely a smoker in this group.

EXPLORE ATTRIBUTES AND CLASS

```
[68]: #Scatterplot Matrix
      cmap = sns.diverging_palette(220, 10, as_cmap=True)

      df_clean_jitter = df_clean[['cardio', 'ap_hi', 'ap_lo', 'cholesterol', 'weight', 'age'
      ↪]]
      df_clean_jitter[['cholesterol', 'age', 'weight']] = df_clean_jitter[['cholesterol',
      ↪'age', 'weight']].values + np.random.rand(len(df_clean_jitter),3)/2
      sns.pairplot(df_clean_jitter, hue="cardio", height=2)

[68]: <seaborn.axisgrid.PairGrid at 0x20f0f11c438>
```

The above pairwise scatter plot depicts the distributions of each of those that had cardiovascular disease and those that did not for each attribute across the center diagonal. As shown in those distributions, there are apparent differences in those that had cardiovascular disease and those that did not. For example, as will be discussed in more detail below, the distributions show that as blood pressure, cholesterol, weight and age increase, the number of those that had cardiovascular disease also increases. In addition, the pairwise scatter plot shows scatterplot distributions between each of the attributes. Some separation between those that had cardiovascular disease and those that did not can be seen in the plots between weight and age, cholesterol and weight, and age and blood pressure, despite the presence of some values seemingly randomly distributed. Therefore, these attributes may be indicative of cardiovascular disease.


```
[69]: f, ax = plt.subplots(figsize=(10, 10))

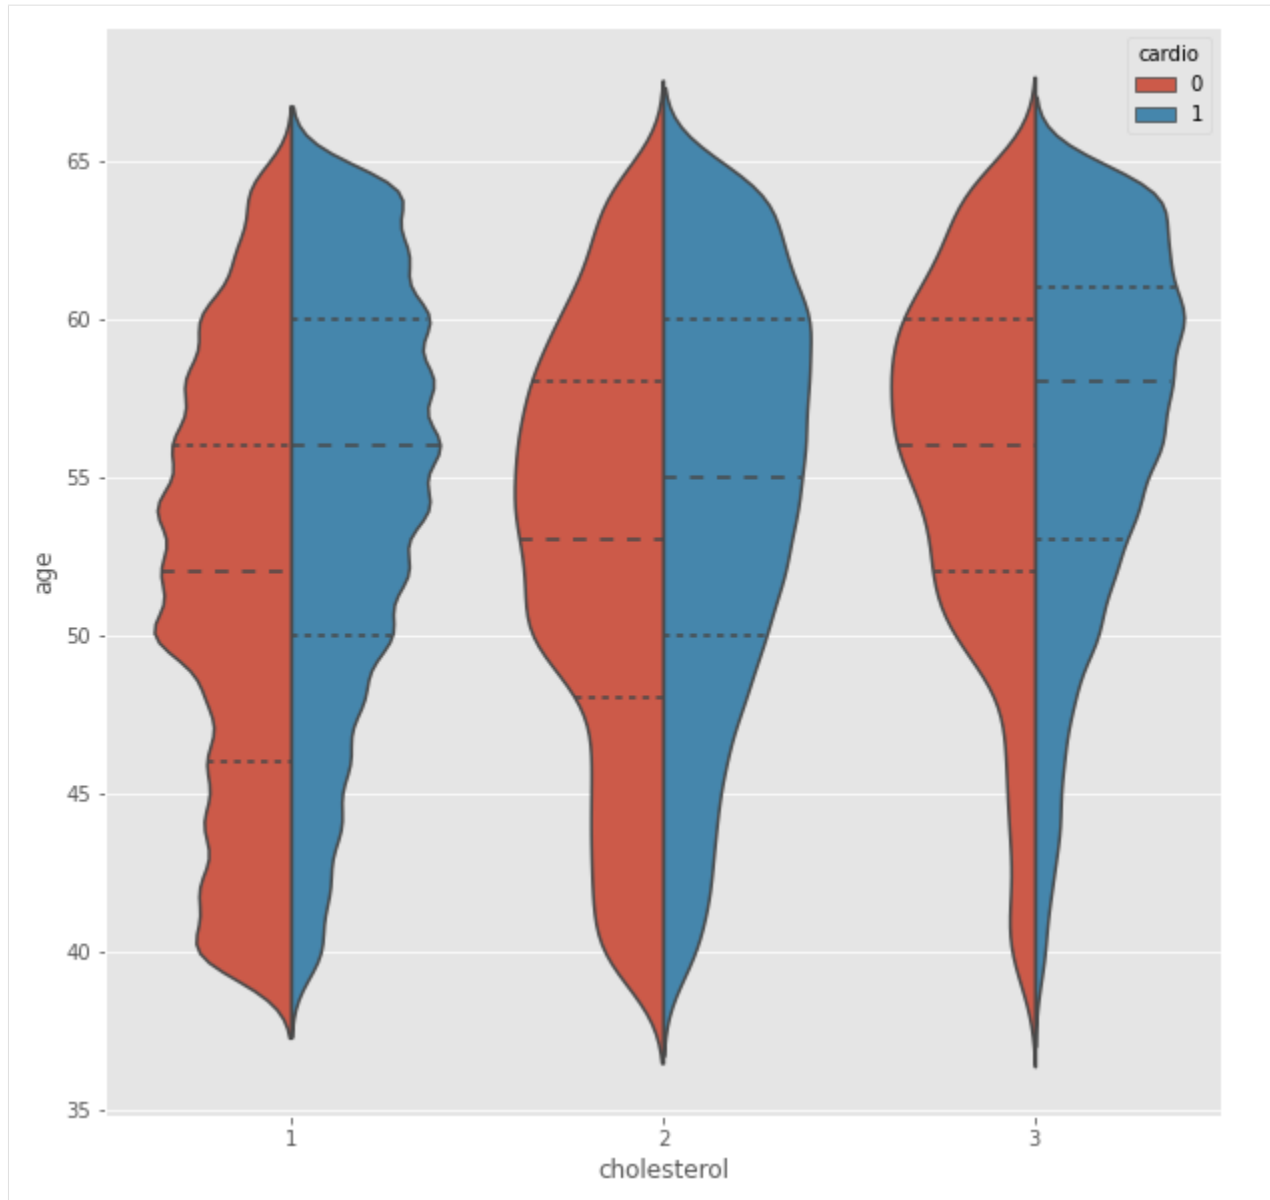
      sns.violinplot(x="cholesterol", y="weight", hue="cardio", data=df_clean,
      split=True, inner="quart")

      f, ax = plt.subplots(figsize=(10, 10))

      sns.violinplot(x="cholesterol", y="age", hue="cardio", data=df_clean,
      split=True, inner="quart")

[69]: <matplotlib.axes._subplots.AxesSubplot at 0x20f0ebe45f8>
```





The above violin plots show the relationship between cholesterol and weight between groups having cardiovascular disease (1), and not having cardiovascular disease. The violin plots show that as the level of cholesterol increases from 1-3, the median difference in weight between the cardiovascular and non-cardiovascular disease groups also increases. It is also apparent that the median weight for cholesterol levels 2-3 for each group represents an increase from cholesterol level 1. This is consistent with the understanding that as cholesterol levels increase, weight also increases.

The violin plots depicting cholesterol and age shows that the distribution between those that had cardiovascular disease and those that did not is consistent across cholesterol levels 2 and 3. For example, the difference between the median ages of those that had cardiovascular disease and those that did not for cholesterol levels 2-3 is approximately equal. This is consistent with that people typically increase cholesterol levels with age as activities such as exercise decrease. However, in the distributions between those that had cardiovascular disease and those that did not for cholesterol level 1 shows that age is a potential indicator of cardiovascular disease. The distribution that had cardiovascular disease at cholesterol level 1 is approximately 56, which is potentially, significantly greater than the median age of those that did not have cardiovascular disease at 52.


```
[70]: f, ax = plt.subplots(figsize=(10, 10))


sns.violinplot(x="cholesterol", y="ap_hi", hue="cardio", data=df_clean,
               split=True, inner="quart")

f, ax = plt.subplots(figsize=(10, 10))

sns.violinplot(x="cholesterol", y="ap_lo", hue="cardio", data=df_clean,
               split=True, inner="quart")

[70]: <matplotlib.axes._subplots.AxesSubplot at 0x20f0f5200b8>
```





Violin plots of blood pressure and cholesterol between those that had and did not have cardiovascular disease appears to include multiple peaks. Otherwise, it appears that those with lower overall blood pressure, between low and high pressure values, suffered less cardiovascular disease. As can be seen, as both low and high blood pressure values increase, the distribution of those with cardiovascular disease also increases. Therefore, it appears as if blood pressure is an indicator of cardiovascular disease.

```
[71]: import plotly.express as px

df_sub2 = df_clean[['age', 'weight', 'cholesterol', 'ap_hi', 'cardio', 'ap_lo']]
#Normalizing the data
df_normalized2 = (df_sub2-df_sub2.mean())/(df_sub2.std())
df_normalized2.cardio = df_sub2.cardio
df_normalized2.cholesterol = df_normalized2.cholesterol+np.random.rand(*df_
↳normalized2.cholesterol.shape)/2
df_normalized2.ap_hi = df_normalized2.ap_hi+np.random.rand(*df_normalized2.ap_hi.
↳shape)/2
```

(continues on next page)

(continued from previous page)

```
fig = px.parallel_coordinates(df_sub2, color='cardio', labels={"age": "Age", "weight":
→ "Weight", "cholesterol": "Cholesterol", "ap_hi": "High Blood Pressure Value",
→ "cardio": "Cardiovascular Disease", "ap_lo": "Low Blood Pressure Value"},
color_continuous_scale=px.colors.diverging.Tealrose,
color_continuous_midpoint=1)
fig.show()
```

Data type cannot be displayed: application/vnd.plotly.v1+json

The above parallel coordinates plot shows interconnections between those that had cardiovascular disease and those that did not, and the attributes that appear to be indicators of cardiovascular disease. Here, those with cardiovascular disease are shown in beige, and those without cardiovascular disease is shown in teal. In the above plot, separation can be seen between age and weight around 43, suggesting that age and weight may interact in predicting cardiovascular disease. Some separation between weight, cholesterol and high blood pressure indicating that as weight, cholesterol and high blood pressure values increase, cardiovascular disease also increases. Lastly, as shown, there is more influence of cardiovascular disease as all of the attributes increase.


```
[72]: from pandas.plotting import parallel_coordinates

#df_sub2 = df_clean[['age', 'weight', 'cholesterol', 'ap_hi', 'cardio', 'ap_lo']]

df_sub = df_clean[['age', 'weight', 'cholesterol', 'ap_hi', 'cardio', 'ap_lo']]
df_sub.cardio = df_sub.cardio=='1'
#normalizing values
df_normalized = (df_sub-df_sub.mean())/(df_sub.std())
df_normalized.cardio = df_sub.cardio
df_normalized.cholesterol = df_normalized.cholesterol+np.random.rand(*df_normalized.
→cholesterol.shape)/2
df_normalized.ap_hi = df_normalized.ap_hi+np.random.rand(*df_normalized.ap_hi.shape)/2

parallel_coordinates(df_normalized,'cardio')

plt.show()
```



The above parallel coordinates plot shows the axes of the attributes for the group that includes cardiovascular disease. Eliminating the group that does not have cardiovascular disease shows more detail after normalization. For example, there appears to be age and weight bands that may be indicative of cardiovascular disease such that outside of this band is indicative of not having cardiovascular disease. Lastly, while most of the axes are clustered 2 and -2, more axes extend from the high values of high and low blood pressure. This is indicative of an increase in cardiovascular disease as the values for high and low blood pressure increase, which also fits the domain knowledge of the problem.

NEW FEATURES

We can combine the height and weight features into a single feature, Body Mass Index. BMI can be calculated using the formula below. In general, we can see a higher instance of cardiovascular disease in patients with a higher BMI.

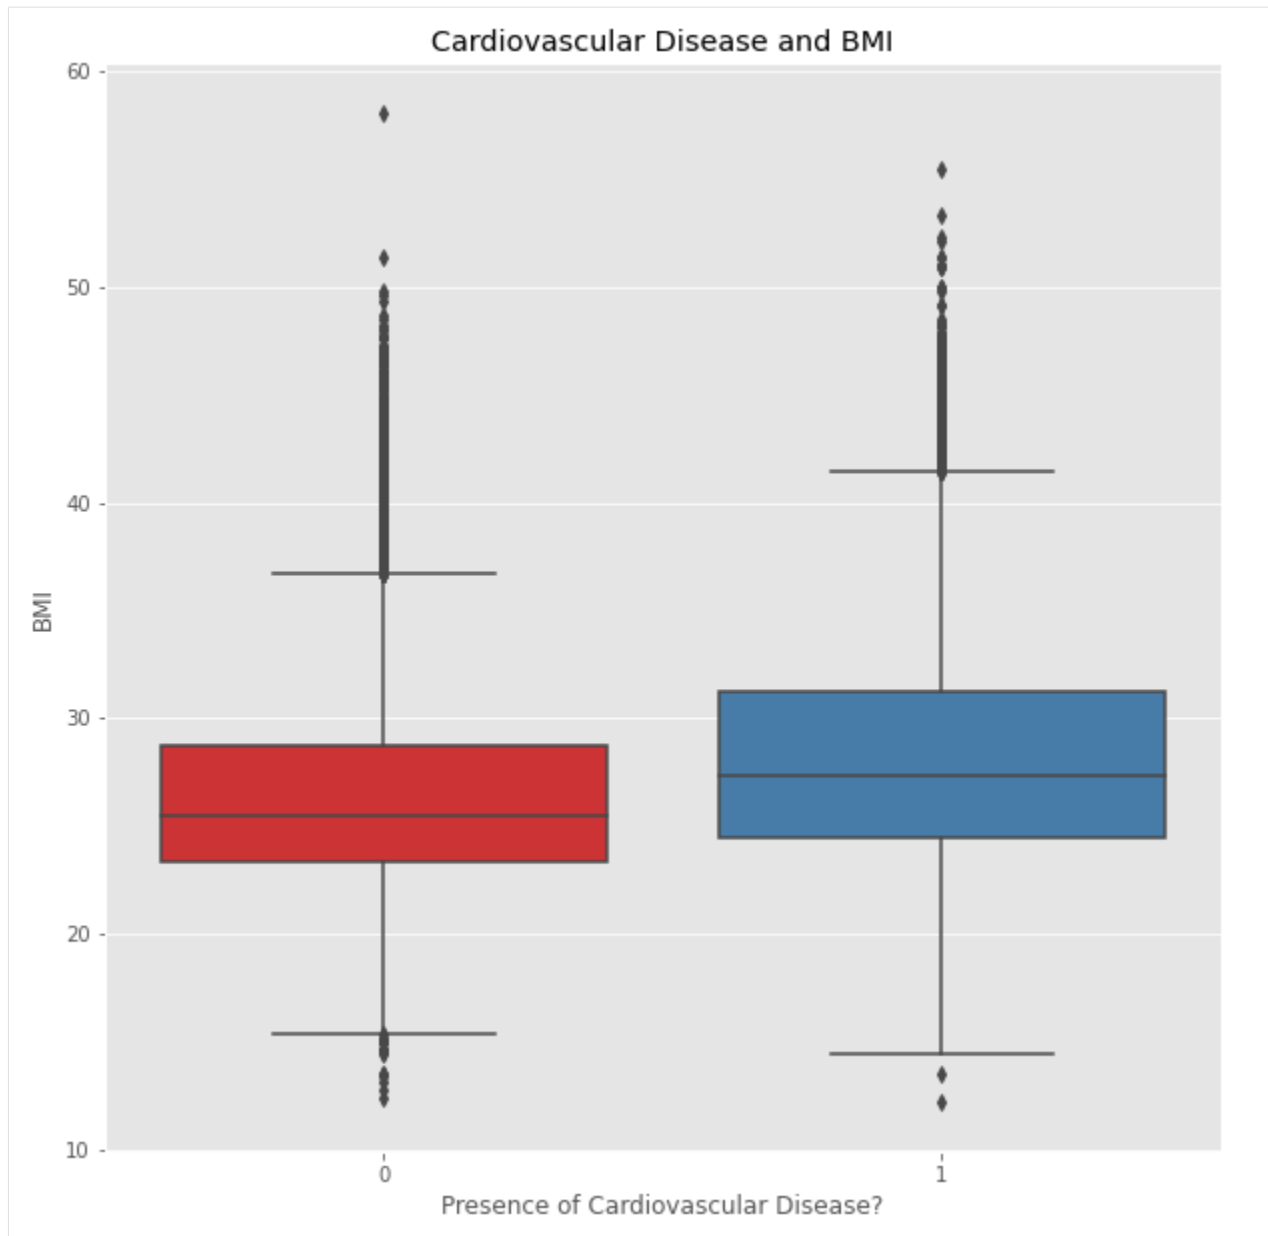
$$BMI = Weight (kg) / Height (m^2)$$

```
[73]: plt.figure(figsize=(10,10))

df_clean['bmi'] = df_clean['weight'] / (df_clean['height']/100)**2

sns.boxplot(x='cardio',
            y='bmi',
            data=df_clean,
            palette="Set1").set(title = 'Cardiovascular Disease and BMI',
                               xlabel = 'Presence of Cardiovascular Disease?',
                               ylabel = 'BMI')

plt.show()
```



For entertainment, we can use some of what we learned in stats so that we don't forget, and do a ttest between the diseased and healthy group. It shows what would expect, that there is some statistical significance in the mean difference of BMI between the two groups (healthy vs diseased).

```
[74]: from scipy import stats
import researchpy as rp

# Let's create 2 sets, one for disease, and another for healthy
disease = df_clean[df_clean['cardio'] == 1]
disease.reset_index(inplace = True)

healthy = df_clean[df_clean['cardio'] == 0]
disease.reset_index(inplace = True)
```

(continues on next page)

(continued from previous page)

```

var='bmi'

# diff = disease[var] - healthy[var]
# stats.probplot(diff, plot= plt)
# plt.title('BMI P-P Plot')
# stats.ttest_ind(disease[var], healthy[var]))

descriptives, results = rp.ttest(disease[var], healthy[var], equal_variances=False)
results

```

```

[74]:
          Welch's t-test      results
0  Difference (bmi - bmi) =      1.8505
1      Degrees of freedom =    69185.9581
2              t =      51.2621
3  Two side test p value =      0.0000
4  Difference < 0 p value =      1.0000
5  Difference > 0 p value =      0.0000
6      Cohen's d =      0.3876
7      Hedge's g =      0.3876
8      Glass's delta =      0.3686
9              r =      0.1913

```

We could also create a BMI category to represent the following four cases. [Source](#)

- Underweight: BMI is less than 18.5.
- Normal weight: BMI is 18.5 to 24.9.
- Overweight: BMI is 25 to 29.9.
- Obese: BMI is 30 or more.

As expected, we can see from that a higher BMI group correlates with a higher chance of being diagnosed with cardiovascular disease.

```

[75]: df_clean['bmiGrp'] = np.where((df_clean.bmi < 18.5), 1, 0)
df_clean['bmiGrp'] = np.where((df_clean.bmi >= 18.5) & (df_clean.bmi < 25), 2, df_
    ↪ clean.bmiGrp)
df_clean['bmiGrp'] = np.where((df_clean.bmi >= 25) & (df_clean.bmi < 30), 3, df_clean.
    ↪ bmiGrp)
df_clean['bmiGrp'] = np.where((df_clean.bmi >= 30), 4, df_clean.bmiGrp)

df_grouped = df_clean.groupby(by=['bmiGrp'])
print ("Percentage of Caridovascular Disease in each BMI group:")
print (df_grouped.cardio.sum() / df_grouped.cardio.count() *100)

Percentage of Caridovascular Disease in each BMI group:
bmiGrp
1      26.917058
2      40.545001
3      51.286476
4      62.745980
Name: cardio, dtype: float64

```

```

[76]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15,8))

sns.boxplot(x='cardio', y='bmiGrp', data=df_clean, ax=ax1)
#sns.barplot(x='bp', y='cardio', data=df, ax=ax2)


```

(continues on next page)

(continued from previous page)

```
cardio = pd.crosstab([df_clean['bmiGrp']], df_clean['cardio'].astype(bool))
cardo_rate = cardio.div(cardio.sum(1).astype(float), axis=0) # normalize the value
cardo_rate.plot(kind='barh', stacked=True, ax=ax2)
```

```
plt.show()
```



We can also create a new feature to categorize the Blood Pressure values. [Source](#)

Blood Pressure Categories



BLOOD PRESSURE CATEGORY	SYSTOLIC mm Hg (upper number)		DIASTOLIC mm Hg (lower number)
NORMAL	LESS THAN 120	and	LESS THAN 80
ELEVATED	120 – 129	and	LESS THAN 80
HIGH BLOOD PRESSURE (HYPERTENSION) STAGE 1	130 – 139	or	80 – 89
HIGH BLOOD PRESSURE (HYPERTENSION) STAGE 2	140 OR HIGHER	or	90 OR HIGHER
HYPERTENSIVE CRISIS (consult your doctor immediately)	HIGHER THAN 180	and/or	HIGHER THAN 120

©American Heart Association

heart.org/bplevels

As with BMI Categories, we can see that a higher BP group corresponds to a higher chance of being diagnosed with

cardiovascular disease.

```
[77]: # Create blood pressure categories

df_clean['bp'] = np.where((df_clean.ap_hi < 120) & (df_clean.ap_lo < 80), 1, 0)
df_clean['bp'] = np.where((df_clean.ap_hi >= 120) & (df_clean.ap_hi < 130) & (df_
↳ clean.ap_lo < 80), 2, df_clean.bp)
df_clean['bp'] = np.where((df_clean.ap_hi >= 130) & (df_clean.ap_hi < 140) | ((df_
↳ clean.ap_lo >= 80) & (df_clean.ap_lo < 90)), 3, df_clean.bp)
df_clean['bp'] = np.where((df_clean.ap_hi >= 140) | (df_clean.ap_lo >= 90), 4, df_
↳ clean.bp)
df_clean['bp'] = np.where((df_clean.ap_hi > 180) | (df_clean.ap_lo > 120), 5, df_
↳ clean.bp)

df_grouped = df_clean.groupby(by=['bp'])
print ("Percentage of cardio disease in each Blood Pressure group:")
print (df_grouped.cardio.sum() / df_grouped.cardio.count() *100)


Percentage of cardio disease in each Blood Pressure group:
bp
1      22.155939
2      32.253968
3      38.524389
4      78.158972
5      85.250000
Name: cardio, dtype: float64
```

```
[78]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15,8))

sns.boxplot(x='cardio', y='bp', data=df_clean, ax=ax1)
#sns.barplot(x='bp', y='cardio', data=df, ax=ax2)

cardio = pd.crosstab([df_clean['bp']],df_clean.cardio.astype(bool))
cardo_rate = cardio.div(cardio.sum(1).astype(float), axis=0) # normalize the value
cardo_rate.plot(kind='barh', stacked=True, ax=ax2)

plt.show()
```



Let's create a new feature correlation heatmap, but instead use the new features we created. In this case, we can see a better correlation between our groups for BMI and BP with respect to our response variable (cardio). Please note, we used a mask below to set correlations less than .2 to 0, for a better visualization.

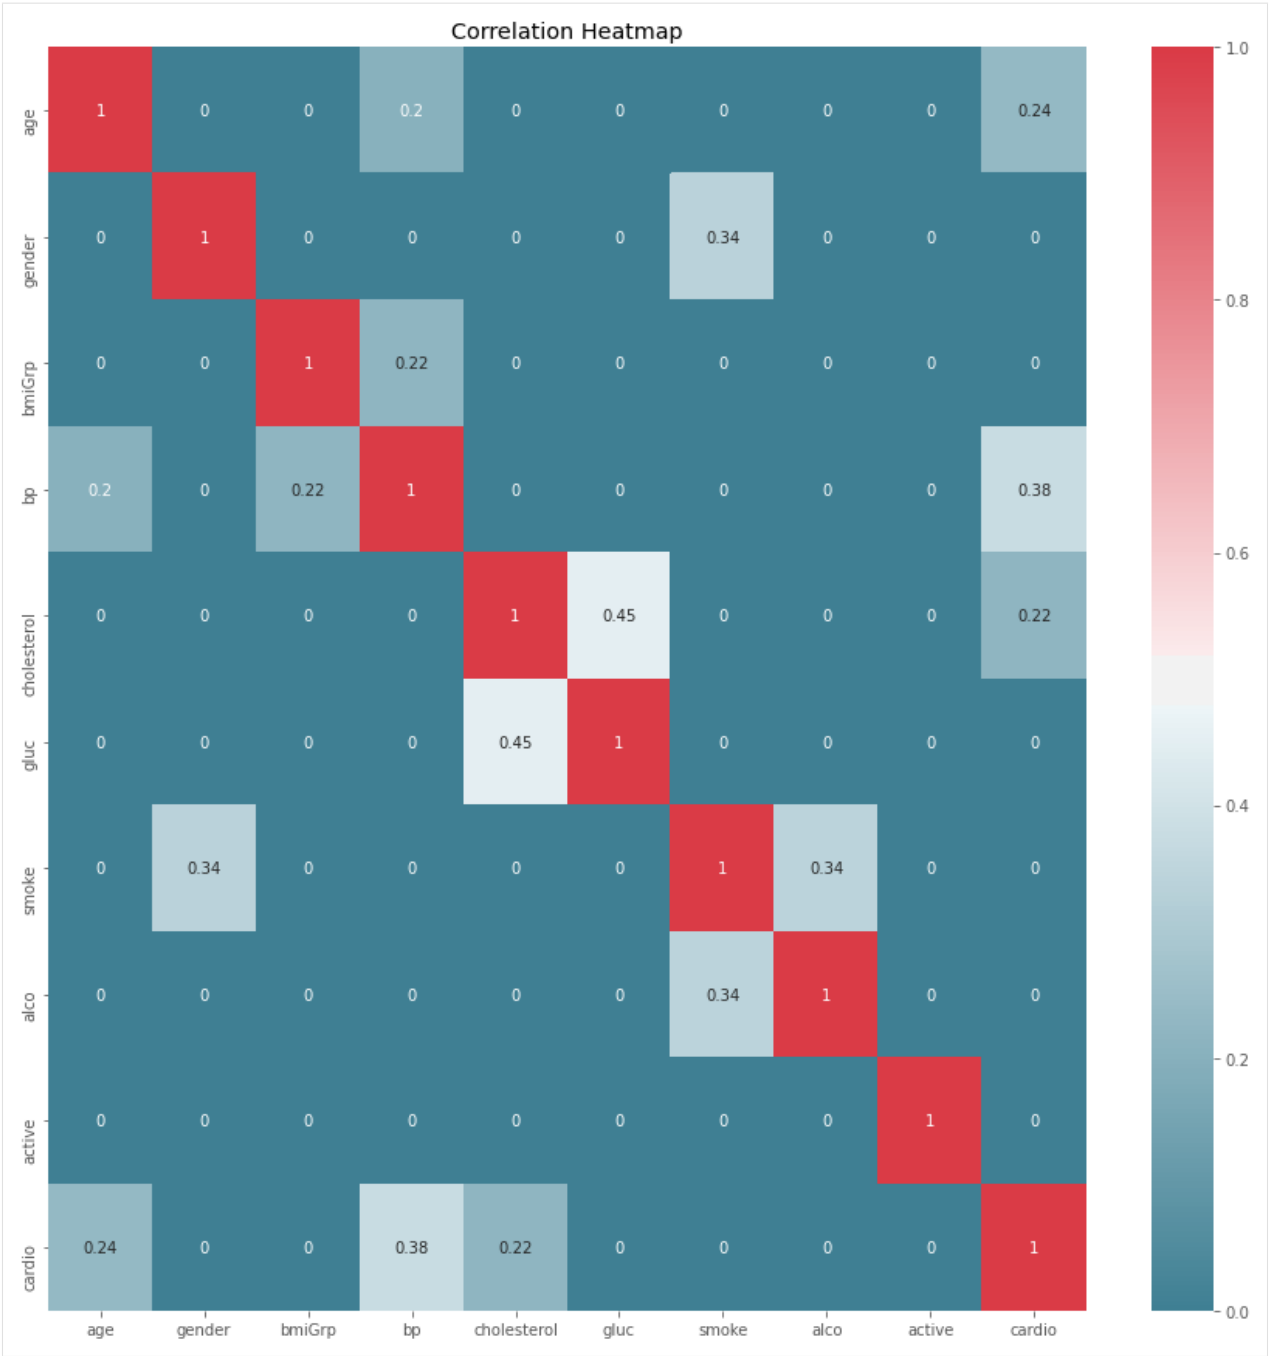
```
[79]: features = ['age', 'gender', 'bmiGrp', 'bp', 'cholesterol', 'gluc', 'smoke', 'alco',
→ 'active', 'cardio']

plt.figure(figsize=(15,15))

# Use an easier to see colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)

correlation = df_clean[features].corr()
correlation[np.abs(correlation)<.2] = 0

sns.heatmap(correlation, annot = True, cmap=cmap).set(title = 'Correlation Heatmap')
plt.show()
```

EXCEPTIONAL WORK

9.1 Store feature matrix as an ndarray

```
[80]: # Full Model
X_cols = ['age', 'gender', 'height', 'weight', 'ap_hi', 'ap_lo', 'cholesterol', 'gluc'
↪', 'smoke', 'alco', 'active']
# *dcrouthamel - test with new features
# X_cols = ['age', 'gender', 'bmiGrp', 'bp', 'cholesterol', 'gluc', 'smoke', 'alco',
↪'active']
X = df_clean[X_cols].to_numpy()

[81]: type(X)

[81]: numpy.ndarray
```

9.2 Store response vector

```
[82]: y = df_clean['cardio'].to_numpy()
```

Full Model consisting of all features with standardized values.

```
[83]: from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_
↪state=0)

sc = StandardScaler()
X_train_std = sc.fit_transform(X_train)
X_test_std = sc.transform(X_test)

# logreg = LogisticRegression()
# logreg.fit(X_train, y_train)
logreg = LogisticRegression()
logreg.fit(X_train_std, y_train)

[83]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='auto', n_jobs=None, penalty='l2',
random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
warm_start=False)
```

```
[84]: y_pred = logreg.predict(X_test_std)
print('Accuracy of the log reg model on the test data: {:.2f}'.format(logreg.score(X_
↪test_std, y_test)))
```

Accuracy of the log reg model on the test data: 0.73

```
[85]: from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(y_test, y_pred)
print(confusion_matrix)
```

```
[[8136 2293]
 [3436 7128]]
```

```
[86]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0		0.70	0.78	10429
1		0.76	0.67	10564
accuracy				0.73 20993
macro avg		0.73	0.73	0.73 20993
weighted avg		0.73	0.73	0.73 20993

```
[87]: from IPython.display import Markdown as md
```

```
[88]: ### Logistic Regression Metrics
```

```
md(f"***True Positives:** {confusion_matrix[1, 1]} \n\n **True Negatives:**
↪{confusion_matrix[0, 0]} \n\n **False Positives:** {confusion_matrix[0, 1]} \n\n
↪**False Negatives:** {confusion_matrix[1, 0]}")
```

[88]: **True Positives:** 7128

True Negatives: 8136

False Positives: 2293

False Negatives: 3436

```
[89]: md(f"***Accuracy:** { format(( confusion_matrix[1, 1] + confusion_matrix[0, 0] ) /
↪confusion_matrix.sum(), '.3f')} \n\n-how often we were correct overall")
```

[89]: **Accuracy:** 0.727

-how often we were correct overall

```
[90]: md(f"***Error:** { format(( confusion_matrix[0, 1] + confusion_matrix[1, 0] ) /
↪confusion_matrix.sum(), '.3f')} \n\n-how often we were incorrect overall")
```

[90]: **Error:** 0.273

-how often we were incorrect overall

```
[91]: md(f"***Sensitivity/ Recall:** { format(( confusion_matrix[1, 1] ) / confusion_
↪matrix[1].sum(axis=0), '.3f')} \n\n-when the patient actually had CVD, how often
↪were we correct")
```

[91]: **Sensitivity/ Recall: 0.675**

-when the patient actually had CVD, how often were we correct

```
[92]: md(f"***Specificity:** { format(( confusion_matrix[0, 0] ) / confusion_matrix[0].
      ↳sum(), '.3f') }\n\n-when the patient did not had CVD, how often were we correct")
```

[92]: **Specificity: 0.780**

-when the patient did not had CVD, how often were we correct

```
[93]: md(f"***False Postive Rate:** { format(( confusion_matrix[0, 1] ) / ( confusion_
      ↳matrix[0, 0] + confusion_matrix[0, 1] ), '.3f') }\n\n-when the patient did not had_
      ↳CVD, how often were we incorrect")
```

[93]: **False Postive Rate: 0.220**

-when the patient did not had CVD, how often were we incorrect

```
[94]: md(f"***Precision:** { format(( confusion_matrix[1, 1] ) / ( confusion_matrix[1, 1] +_
      ↳confusion_matrix[0, 1] ), '.3f') }\n\n-how precise were we when classifying the_
      ↳patient as having CVD")
```

[94]: **Precision: 0.757**

-how precise were we when classifying the patient as having CVD

```
[95]: from sklearn.metrics import plot_confusion_matrix
      from sklearn.metrics import plot_roc_curve


      fig = plt.figure(1, figsize=(20, 5))

      chart_1 = fig.add_subplot(121)
      chart_2 = fig.add_subplot(122)

      plot_confusion_matrix(logreg, X_test_std, y_test, normalize='true', ax=chart_1)
      chart_1.set_title('Confusion Matrix')

      plot_roc_curve(logreg, X_test_std, y_test, ax=chart_2)
      chart_2.set_title('ROC Curve')

      plt.show()
```



Principal Component Analysis is a data compression technique that can reduce the dimensionality of a data set. It does this by finding the maximum variance in a higher dimensional space and project that onto a new space with fewer

dimensions. Although our data set doesn't have a huge number of features, let's explore what PCA can do for us.

First, we'll define a function that can plot two principal components and a decision boundary. This code was taken from chapter 4 of Python Machine Learning, by Vahid Mirjalili and Sebastian Raschka.

```
[96]: # Chapter 4 of book
from matplotlib.colors import ListedColormap
def plot_decision_regions(X, y, classifier, resolution=0.02):

    # setup marker generator and color map
    markers = ('s', 'x', 'o', '^', 'v')
    colors = ('red', 'blue', 'lightgreen', 'gray', 'cyan')
    cmap = ListedColormap(colors[:len(np.unique(y))])

    # plot the decision surface
    x1_min, x1_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    x2_min, x2_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution),
                           np.arange(x2_min, x2_max, resolution))
    Z = classifier.predict(np.array([xx1.ravel(), xx2.ravel()]).T)
    Z = Z.reshape(xx1.shape)
    plt.contourf(xx1, xx2, Z, alpha=0.4, cmap=cmap)
    plt.xlim(xx1.min(), xx1.max())
    plt.ylim(xx2.min(), xx2.max())

    # plot examples by class
    for idx, cl in enumerate(np.unique(y)):
        plt.scatter(x=X[y == cl, 0],
                    y=X[y == cl, 1],
                    alpha=0.6,
                    color=cmap(idx),
                    edgecolor='black',
                    marker=markers[idx],
                    label=cl)
```


Below, we'll create a PCA classifier with two components and fit the reduced data set to a logistic regression model.

```
[97]: from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler

pca = PCA(n_components=2)
# sc = StandardScaler()
# X_train_std = sc.fit_transform(X_train)
# X_test_std = sc.transform(X_test)

# dimensionality reduction:
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# fitting the logistic regression model on the reduced dataset:
logreg = LogisticRegression()
logreg.fit(X_train_pca, y_train)
plot_decision_regions(X_train_pca, y_train, classifier=logreg)
plt.xlabel('PC 1')
plt.ylabel('PC 2')
plt.legend(loc='lower left')
plt.tight_layout()
plt.show()
```



```
[98]: pca.explained_variance_ratio_
```

```
[98]: array([0.55118299, 0.24041201])
```

Above, we can see that the two components account nearly 80 percent of the variance. The figure though doesn't show a good separation. This could be due to the fact that we need to visualize it in more than 2 dimensions, i.e., run PCA to account for a higher number of components and variance.

Below we output some metrics related to using the reduced dimensionality set in a Logistic Regression model. We see that the accuracy is comparable to the full model created previously. Accuracy is an acceptable metric for a balance data set. However, in the case of medical diagnosis, Recall or Sensitivity is an important metric. It describes the proportion of patients correctly diagnosed with CVD. If this number is low, patients won't be correctly identified and won't receive the treatment they should.

PCA shows .62 for recall, whereas the previous model is .67. Because of that, I would go with the full model.

```
[99]: y_pred = logreg.predict(X_test_pca)
print('Accuracy of the log reg model on the test data: {:.2f}'.format(logreg.score(X_
    test_pca, y_test)))
```

```
Accuracy of the log reg model on the test data: 0.71
```

```
[100]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0		0.68	0.80	10429
1		0.76	0.62	10564
accuracy				0.71 20993
macro avg		0.72	0.71	0.71 20993
weighted avg		0.72	0.71	0.71 20993

```
[101]: from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import plot_roc_curve
```

(continues on next page)

(continued from previous page)

```

fig = plt.figure(1, figsize=(20, 5))


chart_1 = fig.add_subplot(121)
chart_2 = fig.add_subplot(122)

plot_confusion_matrix(logreg, X_test_pca, y_test, normalize='true', ax=chart_1)
chart_1.set_title('Confusion Matrix')

plot_roc_curve(logreg, X_test_pca, y_test, ax=chart_2)
chart_2.set_title('ROC Curve')

plt.show()

```



PCA (unsupervised) does not use the target variable. It attempts to maximize the variance in the feature set. LDA (supervised) can use our target information and maximize the class separability.

Note that LDA will produce $N-1$ components, where N is the number of classes in our target. So in our case, it would produce 1 components.

We can use the below to generate an LDA feature space and train a logistic regression model against it.

```

[102]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
        from sklearn.preprocessing import StandardScaler

        lda = LinearDiscriminantAnalysis()
        # sc = StandardScaler()
        # X_train_std = sc.fit_transform(X_train)
        # X_test_std = sc.transform(X_test)

        X_train_lda = lda.fit_transform(X_train, y_train)
        X_test_lda = lda.transform(X_test)

        # fitting the logistic regression model on the lda dataset:
        logreg = LogisticRegression()
        logreg.fit(X_train_lda, y_train)

[102]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, l1_ratio=None, max_iter=100,
                           multi_class='auto', n_jobs=None, penalty='l2',
                           random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                           warm_start=False)

```



```
[103]: y_pred = logreg.predict(X_test_lda)
print('Accuracy of the log reg model on the test data: {:.2f}'.format(logreg.score(X_
↪test_lda, y_test)))
```

Accuracy of the log reg model on the test data: 0.73

```
[104]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

```
precision    recall  f1-score   support

           0       0.70      0.77      0.74      10429
           1       0.75      0.68      0.71      10564

 accuracy                   0.73      20993
 macro avg              0.73      0.73      0.73      20993
 weighted avg          0.73      0.73      0.73      20993
```

```
[105]: from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import plot_roc_curve


fig = plt.figure(1, figsize=(20, 5))

chart_1 = fig.add_subplot(121)
chart_2 = fig.add_subplot(122)

plot_confusion_matrix(logreg, X_test_lda, y_test, normalize='true', ax=chart_1)
chart_1.set_title('Confusion Matrix')

plot_roc_curve(logreg, X_test_lda, y_test, ax=chart_2)
chart_2.set_title('ROC Curve')

plt.show()
```



Above we can see that our accuracy increased to .73 using LDA, vs .71 with PCA. Additionally, the AUC value is higher and maybe more importantly, the Recall/Sensitivity value increased to .68. I think this helps demonstrate that LDA did a better job since it was able to use our target information.